

What's the Simplest IR System?

- Given a collection of documents and a “free text” query
- How can we get some search results in a simple way?
- *grep-like*: a “sequential scan”
- Simple but ...
 - very inefficient
- Is it effective?



How can we make it more effective AND efficient?